

Austin Tangban

Software Engineer | LLM Product Engineer
aotangban@gmail.com aushtiin.github.io

SUMMARY

Software engineer with **more than half a decade** shipping production systems and two years of applied NLP/LLM work. I build LLM products end to end: human-in-the-loop drafting pipelines and AI-assisted partner matching at Souk, and fine-tuning research at LORIA, where my domain-adapted models read FDA drug labels at **94-96 F1**, about 20 points ahead of the best zero-shot LLM. Equally at home tuning a slow Postgres query and evaluating a model.

EXPERIENCE

SOUK | FULL-STACK ENGINEER

Jan 2026 – Present | Remote (Part-time)

- Souk is an AI platform that finds companies new referral and affiliate partners, runs the outreach, and tracks the deals; I build across the whole product.
- Designed and built end-to-end email infrastructure: automated domain registration (Dynadot API), DNS configuration (Cloudflare API), mailbox provisioning (Dovecot), and outbound sending (Mailgun), fully replacing a paid third-party service.
- Implemented automated domain warmup scheduling and webhook-driven deliverability monitoring, surfaced in a realtime admin dashboard.
- Built AutoWrite, a production LLM pipeline (Anthropic API) that drafts campaign emails and LinkedIn messages with a human-in-the-loop review gate before anything is sent.
- Integrated AI scoring into partner matching, ranking and recommending partners across the network.
- **Tech:** TypeScript, Node.js, Express, React, Bun, BullMQ, Cloudflare, AWS S3, Anthropic API.

PRACTICE LYNCHPYN INC. | FULL-STACK SOFTWARE ENGINEER

Sep 2021 – Dec 2025 | Toronto, Canada (remote)

- Lynchpyn is a coaching platform that leadership teams use to run their planning: goals, KPIs, meetings, and the routines around them.
- Engineered a formula engine that lets teams define custom KPI calculations directly in the app, replacing complex Excel workflows.
- Designed and shipped a routine-templating system powering dynamic team meetings; adopted as the default workflow company-wide.
- Cut average API latency from **120s to under 300ms** through Postgres query optimization and Redis caching layers.
- Built a centralized ETL pipeline consolidating CurveHero, Meta Ads, and Google Ads data into a single Postgres source of truth.
- Automated daily data audits with an LLM job that reasons over records and backfills missing fields.
- **Tech:** Ruby on Rails, React, TypeScript, MobX, Postgres, Redis.

SIRIUSLABS | FULL-STACK BLOCKCHAIN ENGINEER (CONTRACT)

Mar 2021 – Sep 2021 | Delaware, United States (remote)

- Built a blockchain-enabled NFT marketplace: smart contracts, wallet connectivity, and decentralized storage.
- **Tech:** React, Node.js, Solidity, Ethereum, Web3.js, Ethers.js.

PAYQART SOLUTIONS | FULL-STACK ENGINEER (CONTRACT)

Mar 2021 – Aug 2021 | Lagos, Nigeria

- Implemented core features for a Buy-Now-Pay-Later fintech: secure payment workflows, OAuth-based authentication, and real-time transaction monitoring pipelines.
- **Tech:** React, Node.js, Express, MongoDB.

SKILLS

Software Engineering: Python, JavaScript, TypeScript, Ruby, Ruby on Rails, Node.js, Bun, React, Swift, Postgres, MongoDB, Redis, BullMQ, Docker, GraphQL

ML / AI: PyTorch, TensorFlow, HuggingFace Transformers, Scikit-learn, Pandas, NumPy, spaCy, LLM pipeline integration (Anthropic API), human-in-the-loop review systems, extractive QA, text classification, transformer fine-tuning, topic modelling (LDA, BERTopic), model evaluation (F1, Accuracy, BLEU)

NLP AND MACHINE LEARNING PROJECTS

PHARMACOKINETIC DATA EXTRACTION FROM FDA DRUG LABELS LORIA - Nancy, France | NLP

Research Intern, Mar 2026 – Present

- Fine-tuned domain-adapted BERT variants (BioBERT, PubMedBERT) for extractive question answering over FDA drug labels, extracting 8 pharmacokinetic properties at **94–96 F1**.
- Built the document ingestion pipeline from raw PDF to clean text (GROBID, PDFPlumber) with regex and local-LLM cleanup passes.
- Benchmarked zero-shot LLMs against the fine-tuned models; the best (Llama-3.1-8B, 74 F1) trails fine-tuning by about 20 F1 on this task.
- Added a vision-language model stage that recovers PK values existing only in tables, lifting F1 from 46 to 84 on table-borne values.

GEOTHERMAL PROJECT DOCUMENT CLASSIFIER Université de Lorraine

- Built an ML pipeline that classifies French municipal meeting minutes to surface *geothermie* (geothermal energy) projects: PDF ingestion, OCR, cleaning, and document structuring with Python and spaCy.
- Applied long-document transformers (BERT, Longformer) with custom keyword and semantic features.
- Evaluated on imbalanced, domain-specific data; supports energy-policy document retrieval.

POLITICAL DISCOURSE CORPUS CONSTRUCTION AND ANALYSIS (INSIGHT PROJECT) LORIA -

Nancy, France

- Built a structured corpus of 9k messages from raw Discord screenshots via an LLM-assisted extraction pipeline recovering message text, timestamps, usernames, and nested replies.
- Implemented French/English preprocessing with spaCy (language detection, custom stopwords, lemmatization) and metadata enrichment for downstream discourse analysis.
- Applied topic modelling (LDA, BERTopic) and multilingual sentiment analysis (XLM-RoBERTa) to identify recurring themes and tonal patterns across channels and time.

SONG-LYRIC EMOTION CLASSIFIER Université de Lorraine

- Co-developed a labeled corpus of song lyrics across six emotion categories.
- Fine-tuned BERT and DistilBERT (HuggingFace) against classical baselines (SVM, Logistic Regression) with reproducible PyTorch and scikit-learn training pipelines.
- Improved F1 significantly through data cleaning and class rebalancing.

EDUCATION

UNIVERSITÉ DE LORRAINE MASTER OF SCIENCE IN NATURAL LANGUAGE PROCESSING

Nancy, France

UNIVERSITY OF PORT HARCOURT BSc IN PHYSICS/ELECTRONICS

Port Harcourt, Nigeria

LINKS

Website:// aushtiin.github.io Github:// [Aushtiin](#) LinkedIn:// [Austin Tangban](#)